

```
In [1]: # %% 0. Imports & runtime setup
import os, requests, types, numpy as np, pandas as pd, tensorflow as tf
from sklearn.utils import class_weight
from tensorflow.keras import layers, mixed_precision

# Enable mixed precision on any detected GPU (handy for RTX 5090)
if tf.config.list_physical_devices("GPU"):
    mixed_precision.set_global_policy("mixed_float16")

# Simple args object for notebook use
Args = types.SimpleNamespace
args = Args(
    csv="all_countries_videos_cleaned.csv", # path to your metadata CSV
    img_dir="thumbnails",                  # folder for thumbnails
    model="efficientnetv2l",                # "resnet50" or "efficientne
    batch=64,
    epochs=5,
    unfreeze=10,                           # fine-tune last N layers
)
print(args)
```

```
2025-06-06 00:54:14.290485: I tensorflow/core/util/port.cc:153] oneDNN cus
tom operations are on. You may see slightly different numerical results du
e to floating-point round-off errors from different computation orders. To
turn them off, set the environment variable `TF_ENABLE_ONEDNN_OPTS=0`.
2025-06-06 00:54:14.305872: E external/local_xla/xla/stream_executor/cuda/
cuda_fft.cc:477] Unable to register cuFFT factory: Attempting to register
factory for plugin cuFFT when one has already been registered
WARNING: All log messages before absl::InitializeLog() is called are writt
en to STDERR
E0000 00:00:1749196454.320765      2339 cuda_dnn.cc:8310] Unable to register
cuDNN factory: Attempting to register factory for plugin cuDNN when one ha
s already been registered
E0000 00:00:1749196454.325070      2339 cuda_blas.cc:1418] Unable to regist
er cuBLAS factory: Attempting to register factory for plugin cuBLAS when on
e has already been registered
2025-06-06 00:54:14.340418: I tensorflow/core/platform/cpu_feature_guard.c
c:210] This TensorFlow binary is optimized to use available CPU instructio
ns in performance-critical operations.
To enable the following instructions: AVX2 AVX_VNNI FMA, in other operatio
ns, rebuild TensorFlow with the appropriate compiler flags.
namespace(csv='all_countries_videos_cleaned.csv', img_dir='thumbnails', mo
del='efficientnetv2l', batch=64, epochs=5, unfreeze=10)
```

```
In [2]: # %% 1. Backbone registry (swap light vs heavy CNN)
from tensorflow.keras.applications import ResNet50, EfficientNetV2L
from tensorflow.keras.applications.resnet50 import preprocess_input as rn
from tensorflow.keras.applications.efficientnet_v2 import preprocess_inpu

MODELS = {
    "resnet50": {
        "cls": ResNet50,
```

```

        "img_size": (224, 224),
        "preprocess": rn_preprocess,
    },
    "efficientnetv2l": {
        "cls": EfficientNetV2L,
        "img_size": (384, 384),
        "preprocess": ev2_preprocess,
    },
}
conf = MODELS[args.model]
IMG_SIZE = conf["img_size"]

```

```

In [3]: # %% 1. Read CSV and prepare image paths + label
df = pd.read_csv(args.csv)

# Create binary label: 1 if viral, 0 otherwise
df["is_viral"] = ((df["views"] > 100_000) & (df["likes"] > 10_000)).astype(int)

# Construct local image path using video_id (e.g., thumbnails/abc123.jpg)
df["thumbnail_path"] = df["video_id"].apply(lambda vid: f"{args.img_dir}/{vid}.jpg")

# Remove rows where the image file doesn't exist
df = df[df["thumbnail_path"].apply(os.path.exists)].reset_index(drop=True)

print("✅ Valid samples:", len(df))

```

✅ Valid samples: 254606

```

In [4]: # %% 3. tf.data pipeline (image branch + TF-IDF title)
AUTOTUNE = tf.data.AUTOTUNE
VOCAB = 2000

# Build TextVectorization layer using TF-IDF
vectorize = layers.TextVectorization(
    max_tokens=VOCAB,
    output_mode="tf_idf"
)
vectorize.adapt(df["title"])

# Image loading and preprocessing
def load_image(path):
    img = tf.io.read_file(path)
    img = tf.image.decode_jpeg(img, channels=3)
    img = tf.image.resize(img, IMG_SIZE)
    img = conf["preprocess"](img)
    return img

# Wrap inputs into a dictionary
def pack(title, img_path, label):
    return {"title": title, "thumbnail": img_path}, label

# Apply preprocessing to both title and image
def preprocess(inputs, label):
    title = vectorize(inputs["title"])

```

```

        thumbnail = load_image(inputs["thumbnail"])
        label = tf.cast(label, tf.float32)
        label = tf.expand_dims(label, -1)
        return {"title": title, "thumbnail": thumbnail}, label

# Build the full dataset
ds = (
    tf.data.Dataset.from_tensor_slices((df["title"].values, df["thumbnail"]
    .map(pack, num_parallel_calls=AUTOTUNE)
    .map(preprocess, num_parallel_calls=AUTOTUNE)
    .shuffle(2048)
    .batch(args.batch)
    .prefetch(AUTOTUNE)
    )
)

```

```

I0000 00:00:1749196827.731497    2339 gpu_device.cc:2022] Created device /
job:localhost/replica:0/task:0/device:GPU:0 with 13499 MB memory:  -> devi
ce: 0, name: NVIDIA GeForce RTX 4080 SUPER, pci bus id: 0000:01:00.0, comp
ute capability: 8.9

```

```

In [5]: # %% 4. Build multimodal model
# Text branch
text_in = tf.keras.Input(shape=(VOCAB,), name="title")
text_x = layers.Dense(256, activation="relu")(text_in)
text_x = layers.Dropout(0.3)(text_x)

# Image branch
img_in = tf.keras.Input(shape=(*IMG_SIZE, 3), name="thumbnail")
backbone = conf["cls"](include_top=False, weights="imagenet", pooling="av
backbone.trainable = False      # phase 1 frozen
img_x = backbone(img_in, training=False)

# Fusion & output
fusion = layers.concatenate([text_x, img_x])
fusion = layers.Dense(128, activation="relu")(fusion)
fusion = layers.Dropout(0.3)(fusion)
out = layers.Dense(1, activation="sigmoid", dtype="float32")(fusion)

model = tf.keras.Model(inputs={"title": text_in, "thumbnail": img_in}, ou
model.summary()

```

**Model: "functional"**

Layer (type)	Output Shape	Param #	Connected to
title (InputLayer)	(None, 2000)	0	–
cast (Cast)	(None, 2000)	0	title[0][0]
dense (Dense)	(None, 256)	512,256	cast[0][0]
thumbnail (InputLayer)	(None, 384, 384, 3)	0	–
dropout (Dropout)	(None, 256)	0	dense[0][0]
efficientnetv2-l (Functional)	(None, 1280)	117,746,8...	thumbnail[0][0]
concatenate (Concatenate)	(None, 1536)	0	dropout[0][0] efficientnetv2-l[0][0]
dense_1 (Dense)	(None, 128)	196,736	concatenate[0][0]
dropout_1 (Dropout)	(None, 128)	0	dense_1[0][0]
cast_2 (Cast)	(None, 128)	0	dropout_1[0][0]
dense_2 (Dense)	(None, 1)	129	cast_2[0][0]

**Total params:** 118,455,969 (451.87 MB)

**Trainable params:** 709,121 (2.71 MB)

**Non-trainable params:** 117,746,848 (449.17 MB)

```
In [6]: # %% 5. Phase 1 training (only top layers)
labels_np = df.is_viral.values
cw = class_weight.compute_class_weight("balanced", classes=np.unique(labels_np),
class_weights = {i: w for i, w in enumerate(cw)})


model.compile(optimizer="adam", loss="binary_crossentropy", metrics=["accuracy"])
history1 = model.fit(ds, epochs=args.epochs, class_weight=class_weights)
```

Epoch 1/5

```

WARNING: All log messages before absl::InitializeLog() is called are written to STDERR
I0000 00:00:1749197058.668381    2586 service.cc:148] XLA service 0x7f40a8dc2f90 initialized for platform CUDA (this does not guarantee that XLA will be used). Devices:
I0000 00:00:1749197058.668942    2586 service.cc:156] StreamExecutor device (0): NVIDIA GeForce RTX 4080 SUPER, Compute Capability 8.9
2025-06-06 01:04:20.357010: I tensorflow/compiler/mlir/tensorflow/utils/dump_mlir_util.cc:268] disabling MLIR crash reproducer, set env var `MLIR_CRASH_REPRODUCER_DIRECTORY` to enable.
I0000 00:00:1749197063.915557    2586 cuda_dnn.cc:529] Loaded cuDNN version 90300
2025-06-06 01:04:33.789684: E external/local_xla/xla/service/slow_operation_alarm.cc:65] Trying algorithm eng0{} for conv (f16[64,96,96,256]{3,2,1,0}, u8[0]{0}) custom-call(f16[64,96,96,64]{3,2,1,0}, f16[256,3,3,64]{3,2,1,0}), window={size=3x3 pad=1_1x1_1, dim_labels=b01f_o01i->b01f, custom_call_target="__cudnn$convForward", backend_config={"cudnn_conv_backend_config":{"activation_mode":"kNone","conv_result_scale":1,"leakyrelu_alpha":0,"side_input_scale":0},"force_earliest_schedule":false,"operation_queue_id":"0","wait_on_operation_queues":[]}} is taking a while...
2025-06-06 01:04:34.101704: E external/local_xla/xla/service/slow_operation_alarm.cc:133] The operation took 1.314447356s
Trying algorithm eng0{} for conv (f16[64,96,96,256]{3,2,1,0}, u8[0]{0}) custom-call(f16[64,96,96,64]{3,2,1,0}, f16[256,3,3,64]{3,2,1,0}), window={size=3x3 pad=1_1x1_1, dim_labels=b01f_o01i->b01f, custom_call_target="__cudnn$convForward", backend_config={"cudnn_conv_backend_config":{"activation_mode":"kNone","conv_result_scale":1,"leakyrelu_alpha":0,"side_input_scale":0},"force_earliest_schedule":false,"operation_queue_id":"0","wait_on_operation_queues":[]}} is taking a while...
I0000 00:00:1749197090.877400    2586 device_compiler.h:188] Compiled cluster using XLA! This line is logged at most once for the lifetime of the process.


```

**3979/3979**  **609s** 140ms/step - accuracy: 0.8053 - loss: 0.4147  
Epoch 2/5


```

2025-06-06 01:14:08.533122: W tensorflow/core/kernels/data/prefetch_autotuner.cc:52] Prefetch autotuner tried to allocate 113758464 bytes after encountering the first element of size 113758464 bytes.This already causes the autotune ram budget to be exceeded. To stay within the ram budget, either increase the ram budget or reduce element size
2025-06-06 01:14:08.534033: W tensorflow/core/kernels/data/prefetch_autotuner.cc:52] Prefetch autotuner tried to allocate 113758976 bytes after encountering the first element of size 113758976 bytes.This already causes the autotune ram budget to be exceeded. To stay within the ram budget, either increase the ram budget or reduce element size

```


**3979/3979**  **660s** 166ms/step - accuracy: 0.8502 - loss: 0.3361

Epoch 3/5


**1/3979**  **42:50** 646ms/step - accuracy: 0.8281 - loss: 0.3258

2025-06-06 01:25:08.455611: W tensorflow/core/kernels/data/prefetch\_autotuner.cc:52] Prefetch autotuner tried to allocate 113758464 bytes after entering the first element of size 113758464 bytes. This already causes the autotune ram budget to be exceeded. To stay within the ram budget, either increase the ram budget or reduce element size

2025-06-06 01:25:08.455864: W tensorflow/core/kernels/data/prefetch\_autotuner.cc:52] Prefetch autotuner tried to allocate 113758976 bytes after entering the first element of size 113758976 bytes. This already causes the autotune ram budget to be exceeded. To stay within the ram budget, either increase the ram budget or reduce element size


**3979/3979**  **532s** 134ms/step - accuracy: 0.8781 - loss: 0.2822

Epoch 4/5


**1/3979**  **41:24** 625ms/step - accuracy: 0.8281 - loss: 0.3756

2025-06-06 01:34:00.546095: W tensorflow/core/kernels/data/prefetch\_autotuner.cc:52] Prefetch autotuner tried to allocate 113758464 bytes after entering the first element of size 113758464 bytes. This already causes the autotune ram budget to be exceeded. To stay within the ram budget, either increase the ram budget or reduce element size

2025-06-06 01:34:00.546302: W tensorflow/core/kernels/data/prefetch\_autotuner.cc:52] Prefetch autotuner tried to allocate 113758976 bytes after entering the first element of size 113758976 bytes. This already causes the autotune ram budget to be exceeded. To stay within the ram budget, either increase the ram budget or reduce element size


**3979/3979**  **531s** 133ms/step - accuracy: 0.8943 - loss: 0.2484

Epoch 5/5

**1/3979**  **43:02** 649ms/step - accuracy: 0.9219 - loss: 0.2106

2025-06-06 01:42:51.289630: W tensorflow/core/kernels/data/prefetch\_autotuner.cc:52] Prefetch autotuner tried to allocate 113758464 bytes after entering the first element of size 113758464 bytes. This already causes the autotune ram budget to be exceeded. To stay within the ram budget, either increase the ram budget or reduce element size

2025-06-06 01:42:51.289889: W tensorflow/core/kernels/data/prefetch\_autotuner.cc:52] Prefetch autotuner tried to allocate 113758976 bytes after entering the first element of size 113758976 bytes. This already causes the autotune ram budget to be exceeded. To stay within the ram budget, either increase the ram budget or reduce element size

**3979/3979**  **530s** 133ms/step - accuracy: 0.9058 - loss: 0.2231

```
In [7]: # %% 6. Phase 2 fine-tuning & save
        for layer in backbone.layers[-args.unfreeze:]:
            layer.trainable = True

        model.compile(optimizer=tf.keras.optimizers.Adam(1e-5),
                      loss="binary_crossentropy", metrics=["accuracy"])
        history2 = model.fit(ds, epochs=args.epochs, class_weight=class_weights)


        model.save("viral_predictor.h5")
        print("✓ Training complete - model saved to viral_predictor.h5")
```

Epoch 1/5

2025-06-06 01:52:10.800419: W tensorflow/core/kernels/data/prefetch\_autotuner.cc:52] Prefetch autotuner tried to allocate 113758976 bytes after encountering the first element of size 113758976 bytes. This already causes the autotune ram budget to be exceeded. To stay within the ram budget, either increase the ram budget or reduce element size


**3979/3979**  **623s** 144ms/step - accuracy: 0.8791 - loss: 0.2996

Epoch 2/5


**1/3979**  **43:43** 660ms/step - accuracy: 0.9375 - loss: 0.1749

2025-06-06 02:02:08.013509: W tensorflow/core/kernels/data/prefetch\_autotuner.cc:52] Prefetch autotuner tried to allocate 113758464 bytes after encountering the first element of size 113758464 bytes. This already causes the autotune ram budget to be exceeded. To stay within the ram budget, either increase the ram budget or reduce element size

2025-06-06 02:02:08.013780: W tensorflow/core/kernels/data/prefetch\_autotuner.cc:52] Prefetch autotuner tried to allocate 113758976 bytes after encountering the first element of size 113758976 bytes. This already causes the autotune ram budget to be exceeded. To stay within the ram budget, either increase the ram budget or reduce element size


**3979/3979**  **544s** 137ms/step - accuracy: 0.8970 - loss: 0.2422

Epoch 3/5


**1/3979**  **42:51** 646ms/step - accuracy: 0.8125 - loss: 0.3830

2025-06-06 02:11:11.686930: W tensorflow/core/kernels/data/prefetch\_autotuner.cc:52] Prefetch autotuner tried to allocate 113758464 bytes after encountering the first element of size 113758464 bytes. This already causes the autotune ram budget to be exceeded. To stay within the ram budget, either increase the ram budget or reduce element size

2025-06-06 02:11:11.687113: W tensorflow/core/kernels/data/prefetch\_autotuner.cc:52] Prefetch autotuner tried to allocate 113758976 bytes after encountering the first element of size 113758976 bytes. This already causes the autotune ram budget to be exceeded. To stay within the ram budget, either increase the ram budget or reduce element size

**3979/3979**  **536s** 135ms/step - accuracy: 0.9032 - loss: 0.2275



Epoch 4/5

**1/3979**  **41:13** 622ms/step - accuracy: 0.8750 - loss: 0.3901

2025-06-06 02:20:07.909571: W tensorflow/core/kernels/data/prefetch\_autotuner.cc:52] Prefetch autotuner tried to allocate 113758464 bytes after encountering the first element of size 113758464 bytes. This already causes the autotune ram budget to be exceeded. To stay within the ram budget, either increase the ram budget or reduce element size


2025-06-06 02:20:07.909840: W tensorflow/core/kernels/data/prefetch\_autotuner.cc:52] Prefetch autotuner tried to allocate 113758976 bytes after encountering the first element of size 113758976 bytes. This already causes the autotune ram budget to be exceeded. To stay within the ram budget, either increase the ram budget or reduce element size



**3979/3979**  **532s** 133ms/step - accuracy: 0.9080 - loss: 0.2186  
Epoch 5/5  
**1/3979**  **42:20** 639ms/step - accuracy: 0.8438 - loss: 0.4214

2025-06-06 02:28:59.486187: W tensorflow/core/kernels/data/prefetch\_autotuner.cc:52] Prefetch autotuner tried to allocate 113758464 bytes after encountering the first element of size 113758464 bytes. This already causes the autotune ram budget to be exceeded. To stay within the ram budget, either increase the ram budget or reduce element size

2025-06-06 02:28:59.486450: W tensorflow/core/kernels/data/prefetch\_autotuner.cc:52] Prefetch autotuner tried to allocate 113758976 bytes after encountering the first element of size 113758976 bytes. This already causes the autotune ram budget to be exceeded. To stay within the ram budget, either increase the ram budget or reduce element size

**3979/3979**  **535s** 134ms/step - accuracy: 0.9109 - loss: 0.2125

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save\_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my\_model.keras')` or `keras.saving.save\_model(model, 'my\_model.keras')`.  
✓ Training complete - model saved to viral\_predictor.h5